

MLOps, Evaluation, and Observability

Hit rate is binary: did a relevant document appear anywhere in the retrieved results? Mean Reciprocal Rank

Faithfulness measures whether each claim in an answer is supported by the retrieved context. Correctness

LLM-as-judge systems are susceptible to length and style bias. A narrow rubric and explicit instructions re

Experiment tracking answers 'which configuration performed best', while trace debugging answers 'why did

Retrieval content is untrusted. Prompt injection can cause instructions in retrieved content to override the s

PII leakage can be mitigated with redaction at ingestion time and by enforcing document-level access rules